

Explanatory Notes for 6.390

Shauntclair Ruiz (Current TA)

Fall 2023

Contents

7 Neural Networks 1.5 - Back-Propagation and Training	2
7.5 Error back-propagation	2
7.5.1 Review: Gradient Descent	2
7.5.2 Review: Gradient Descent with LLCs	3
7.5.3 Review: LLC as Neuron	4
7.5.4 LLC Forward-Pass	5
7.5.5 LLC Back-propagation	6
7.5.6 Summary of neural network gradient descent: a high-level view	8
7.5.7 A two-neuron network: starting backprop	9
7.5.8 Continuing backprop: One more problem	11
7.5.9 Finishing two-neuron backprop	12
7.5.10 Many layers: Doing back-propagation	14
7.5.11 What do these derivatives equal?	16
7.5.12 Activation Derivatives	18
7.5.13 Loss derivatives	20
7.5.14 Many neurons per layer	21
7.5.15 The chain rule: Matrix form	22
7.5.16 How the Chain Rule changes in Matrix form	25
7.5.17 Relevant Derivatives	26
7.6 Training	29
7.6.1 Comments	29
7.6.2 Pseudocode	29

CHAPTER 7

Neural Networks 1.5 - Back-Propagation and Training

7.5 Error back-propagation

We have a complete neural network: a **model** we can use to make predictions or calculations.

Now, our mission is to **improve** this neural network: even if our hypothesis class is good, we still have to **find** the hypotheses that are useful for our problem.

As usual, we will start out with **randomized** values for our weights and biases: this **initial** neural network will not be useful for anything in particular, but that's why we need to improve it.

For such a complex problem, we definitely can't find an explicit solution, like we did for ridge regression. Instead, we will have to rely on **gradient descent**.

Concept 1

Neural networks are typically optimized using **gradient descent**.

We randomize them because otherwise, if our initialization is $w_i = 0$, we get

$$w^T x + w_0 = 0$$

no matter what input x we have.

7.5.1 Review: Gradient Descent

What does it really mean to do gradient descent on our **network**? Let's remind ourselves of how gradient descent works, and then **build** up to a network.

Concept 2

Gradient descent works based on the following reasoning:

- We have a function we want to **minimize**: our loss function \mathcal{L} , which tells us how **badly** we're doing.

- We want to perform "less badly".

~~~~~

- Our main tool for **improving**  $\mathcal{L}$  is to alter  $\theta$  and  $\theta_0$ .

- These are our **parameters**: we're adjusting our model.

- The **gradient** is our main tool:  $\frac{\partial \mathcal{L}}{\partial \theta}$  tells you the direction to **change**  $\theta$  in order to **decrease**  $\mathcal{L}$ .

~~~~~

- We want to **change** θ to **decrease** \mathcal{L} . Thus, we move in the direction of

$$\Delta \theta = -\eta \frac{\partial \mathcal{L}}{\partial \theta} \quad (7.1)$$

- Remember that η is our **step size**: we can take bigger or smaller steps in each direction.

~~~~~

- We take steps  $\Delta \theta$  (and  $\Delta \theta_0$ ) until we are satisfied with  $\mathcal{L}$ , or it **stops** improving.

## 7.5.2 Review: Gradient Descent with LLCs

Let's start with a familiar example: LLCs.

Our LLC model uses the following equations:

We'll use  $w$  instead of  $\theta$ .

$$z(x) = w^T x + w_0 \quad g(z) = \sigma(z) = \frac{1}{1 + e^{-z}} \quad (7.2)$$

$$\mathcal{L}(g, y) = y \log(g) + (1 - y) \log(1 - g) \quad (7.3)$$

Our goal is to minimize  $\mathcal{L}$  by adjusting  $w$  and  $w_0$ .

So, we want

$$\frac{\partial \mathcal{L}}{\partial w} \quad \text{and} \quad \frac{\partial \mathcal{L}}{\partial w_0} \quad (7.4)$$

We did this by using the **chain rule**:

We'll focus on  $w$ , but the same goes for  $w_0$ .

$$\frac{\partial \mathcal{L}}{\partial w} = \overbrace{\frac{\partial \mathcal{L}}{\partial g}}^{\mathcal{L}(g)} \cdot \frac{\partial g}{\partial w} \quad (7.5)$$

We can break it up further using **repeated** chain rules:

$$\frac{\partial \mathcal{L}}{\partial w} = \overbrace{\frac{\partial \mathcal{L}}{\partial g}}^{\mathcal{L}(g)} \cdot \underbrace{\frac{\partial g}{\partial z}}_{g(z)} \cdot \frac{\partial z}{\partial w} \quad (7.6)$$

Plugging in our derivatives, we get:

$$\frac{\partial \mathcal{L}}{\partial w} = \overbrace{\left( \frac{y}{\sigma} - \frac{1-y}{1-\sigma} \right)}^{\partial \mathcal{L} / \partial g} \cdot \overbrace{\sigma(1-\sigma)}^{\partial g / \partial z} \cdot \underbrace{x}_{\partial z / \partial w} \quad (7.7)$$

### Concept 3

The **chain rule** allows us to take the gradient of **nested functions**, where each function is the **input** to the next one.


Another way to say this is that one function **feeds into** the next.

If you aren't familiar with "nested" functions, consider this example:

If you have functions  $f(x)$  and  $g(x)$ , then  $g(f(x))$  is the **nested** combination, where the output of  $f$  is the input of  $g$ .

## 7.5.3 Review: LLC as Neuron

Remember that we can represent our LLC as a **neuron**: this could give us the first idea for how to train our **neural network**!




As usual, our first unit  $\Sigma$  is our **linear** component. The output is  $z$ , nothing different from before with LLC.

The **output** of  $\sigma$ , which we wrote before as  $g$ , is now  $a$ .

Something we neglected before: this diagram is **missing** the **loss function**. Let's create a small unit for that.


$\mathcal{L}(a, y)$  has **two** inputs: our predicted value  $a$ , and the correct value  $y$ .

Remember that  $x$  is a whole vector of values, which we've condensed into one variable.




We have two inputs to our loss function.

We **combine** these into a single unit to get:



Our full unit!



### 7.5.4 LLC Forward-Pass

Now, we can do gradient descent like before. We want to get the effect our **weight** has on our **loss**.

But, this time, we'll pair it with a **visual** that is helpful for understanding how we **train** neural networks.

First, one important consideration:

As we saw above, the **gradient** we get might rely on  $z$ ,  $a$ , or  $\mathcal{L}(a, y)$ . So, before we do anything, we have to **compute** these values.

Each step **depends** on the last: this is what the **forward** arrows represent. We call this a **forward pass** on our neural network.

#### Definition 4

A **forward pass** of a neural network is the process of sending information "**forward**" through the neural network, starting from the **input**.

This means the **input** is fed into the **first** layer, and that output is fed into the **next** layer, and so on, until we reach our **final** result and **loss**.

**Example:** If we had

- $f(x) = x + 2$

- $g(f) = 3f$
- $h(g) = \sin(g)$

Then, a forward pass with the input  $x = 10$  would have us go function-by-function:


- $f(10) = 10 + 2$
- $g(f) = 3 \cdot 12$
- $h(g) = \sin(36)$

So, by "forward", we mean that we apply each function, one after another.

In our case, this means computing  $z$ ,  $a$ , and  $\mathcal{L}(a, y)$ .

### 7.5.5 LLC Back-propagation

Now that we have all of our values, we can get our gradient. Let's **visualize** this process.



We want to link  $\mathcal{L}$  to  $w$ . In order to do that, we need to **connect** each thing in between.

- This lets us **combine** lots of simple **links** to get our more complicated result.


We can also call this "chaining together" lots of derivatives.

Loss  $\mathcal{L}$  is what we really care about. So, what is the loss directly **connected** to? The **activation**,  $a$ .

- Our loss function  $\mathcal{L}(a, y)$  contains information about how  $\mathcal{L}$  is linked to  $a$ .

$$\overbrace{\frac{\partial \mathcal{L}}{\partial a}}^{\text{Loss unit}} \quad (7.8)$$

We send this information backwards, so it can be used later.




Now, we're on the  $\sigma(z)$  unit.

- The  $\sigma(z)$  unit contains information about how  $a$  is linked to  $z$ .
- We've connected  $\mathcal{L}$  to  $a$ , and  $a$  to  $z$ . We chain them together, connecting  $\mathcal{L}$  to  $z$ .

$$\underbrace{\frac{\partial \mathcal{L}}{\partial a}}_{\text{Loss unit}} \cdot \underbrace{\frac{\partial a}{\partial z}}_{\text{Activation function}} \quad (7.9)$$


We haven't reached  $w$  yet, so we send this information further back.



Finally, we reach  $\Sigma$ .

- The  $\Sigma$  unit contains information about how  $z$  is linked to  $w$ .
- Finally, we have a chain of links, that allows us to connect  $\mathcal{L}$  to  $w$ .

This last derivative uses  $x$ , because  $w^T x + w_0 = z$ .





And, we built our chain rule! This contains the **information** of the derivatives from **every** unit.

$$\frac{\partial \mathcal{L}}{\partial w} = \overbrace{\frac{\partial \mathcal{L}}{\partial a}}^{\text{Loss unit}} \cdot \overbrace{\frac{\partial a}{\partial z}}^{\text{Activation}} \cdot \overbrace{\frac{\partial z}{\partial w}}^{\text{Linear subunit}} \quad (7.10)$$

Moving backwards like this is called **back-propagation**.

#### Definition 5

**Back-propagation** is the process of moving "**backwards**" through your network, starting at the **loss** and moving back layer-by-layer, and gathering terms in your **chain rule**.

We call it "**propagation**" because we send backwards the **terms** of our chain rule about later derivatives.

An **earlier** unit (closer to the "left") has all of the **derivatives** that come after (to the "right" of) it, along with its own term.

### 7.5.6 Summary of neural network gradient descent: a high-level view

So, with just this, we have built up the basic idea of how we **train** our model: now that we have the gradient, we can do **gradient descent** like we normally do!

This summary covers some things we haven't fully discussed. We'll continue digging into the topic!

**Concept 6**

We can do **gradient descent** on a **neural network** using the ideas we've built up:


- Do a **forward pass**, where we compute the value of each **unit** in our model, passing the information **forward** - each layer's **output** is the next layer's **input**.
  - We finish by getting the **loss**.
- Do **back-propagation**: build up a **chain rule**, starting at the **loss** function, and get each unit's **derivative** in **reverse order**.
  - **Reverse** order: if you have 3 layers, you want to get the 3rd layer's **derivatives**, then the 2nd layer, then the 1st.
  - **Each weight** vector has its own **gradient**: we'll deal with this later, but we need to calculate one for each of them.
- Use your chain rule to get the **gradient**  $\frac{\partial \mathcal{L}}{\partial w}$  for your **weight** vector(s). Take a **gradient descent** step.
- **Repeat** until satisfied, or your model **converges**.

### 7.5.7 A two-neuron network: starting backprop


Above, we mention "each layer": we'll now transition to a **two-neuron** system, so we have "two layers". Then, we'll build up to many layers.

Remember, though, that the **ideas** represented here are just extensions of what we did above.

Let's get a look at our **two-neuron** system, now with our **loss** unit:




And unpack it:



We want to do **back-propagation** like we did before. This time, we have **two** different layers of weights:  $w^1$  and  $w^2$ . Does this cause any problems?


It turns out, it doesn't! We mentioned in the first part of chapter 7 that we can treat the **output** of the **first** layer  $a^1$  as the same as if it were an **input**  $x$ .

This is one of the biggest benefits of neural network layers!



Now, we can do backprop safely.

"Backprop" is a common shortening of "back-propagation".



We can get:

$$\frac{\partial \mathcal{L}}{\partial w^2} = \underbrace{\frac{\partial \mathcal{L}}{\partial a^2}}_{\text{Loss unit}} \cdot \underbrace{\frac{\partial a^2}{\partial z^2}}_{\text{Activation}} \cdot \underbrace{\frac{\partial z^2}{\partial w^2}}_{\text{Linear}} \quad (7.11)$$

The same format as for our **one-neuron** system! We now have a gradient we can update for our **second** weight vector.

But what about our **first** weight vector?

## 7.5.8 Continuing backprop: One more problem

We need to continue further to reach our **earlier** weights: this is why we have to work **backward**.


### Concept 7

We work **backward** in **back-propagation** because every layer after the **current** one **affects** the gradient.

Our current layer **feeds** into the next layer, which feeds into the layer after that, and so on. So this layer affects **every** later layer, which then affect the loss.

So, to see the effect on the **output**, we have to **start** from the **loss**, and get every layer **between** it and our weight vector.

Remember that when we say "f feeds into g", we mean that the output of f is the input to g.



We have one problem, though:

We just gathered the derivative  $\partial \mathcal{L} / \partial w^2$ . If we wanted to continue the chain rule, we would expect to add more terms, like:

$$\frac{\partial w^2}{\partial a^1} \quad (7.12)$$

The problem is, what is  $w^2$ ? It's a vector of constants.

$$w^2 = \begin{bmatrix} w_1^2 \\ w_2^2 \\ \vdots \\ w_n^2 \end{bmatrix}, \quad \text{Not a function of } a^1! \quad (7.13)$$

Since our current derivative includes  $w^2$ , we would continue it with a  $w^2$  in the "top" of a derivative,

$$\frac{\partial \mathcal{L}}{\partial w^2} \frac{\partial w^2}{\partial r}$$


We're not sure what "r" is yet.

That derivative above is going to be **zero**! In other words,  $w^2$  isn't really the **input** to  $z^2$ : it's a **parameter**.

So, we can't end our derivative with  $w^2$ . Instead, we have to use something else.  $z^2$ 's real input is  $a^1$ , so let's go directly to that!

We were building our chain rule by combining inputs with outputs: that's what links two layers together.

So, it should make sense that using something like  $w$  (that doesn't link two layers) prevents us from making a longer chain rule.



Using this allows us to move from layer 2 to layer 1.

Now, we have our new chain rule:

$$\frac{\partial \mathcal{L}}{\partial a^1} = \overbrace{\frac{\partial \mathcal{L}}{\partial a^2} \cdot \frac{\partial a^2}{\partial z^2}}^{\text{Other terms}} \cdot \overbrace{\frac{\partial z^2}{\partial a^1}}^{\text{Link Layers}} \quad (7.14)$$

### Concept 8

For our **weight gradient** in layer  $l$ , we have to end our **chain rule** with

$$\frac{\partial z^l}{\partial w^l}$$

So we can get

$$\frac{\partial \mathcal{L}}{\partial w^l} = \overbrace{\frac{\partial \mathcal{L}}{\partial z^l}}^{\text{Other terms}} \cdot \overbrace{\frac{\partial z^l}{\partial w^l}}^{\text{Get weight grad}}$$

However, because  $w^l$  is not the **input** of layer  $l$ , we can't use it to find the gradient of **earlier layers**.

Instead, we use

$$\frac{\partial z^l}{\partial a^{l-1}} \quad (7.15)$$

To "**link together**" two different layers  $l$  and  $l - 1$  in a **chain rule**.


## 7.5.9 Finishing two-neuron backprop

Now that we have safely connected our layers, we can do the rest of our gradient. First, let's lump together everything we did before:

In this section, we compressed lots of derivatives into


$$\frac{\partial \mathcal{L}}{\partial z^l}$$

Don't let this alarm you, this just hides our long chain of derivatives!




All the info we need is stored in this derivative: it can be written out using our friendly chain rule from earlier.

Now, we can add our remaining terms. It's the same as before: we want to look at the pre-activation



And finally, our input:



We can get our second chain rule

$$\frac{\partial \mathcal{L}}{\partial w^1} = \overbrace{\frac{\partial \mathcal{L}}{\partial a^1}}^{\text{Other layers}} \cdot \overbrace{\frac{\partial a^1}{\partial z^1} \cdot \frac{\partial z^1}{\partial w^1}}^{\text{Layer 1}} \quad (7.16)$$

Which, in reality, looks much bigger:

$$\frac{\partial \mathcal{L}}{\partial w^1} = \overbrace{\left( \frac{\partial \mathcal{L}}{\partial a^2} \right)}^{\text{Loss unit}} \cdot \overbrace{\left( \frac{\partial a^2}{\partial z^2} \cdot \frac{\partial z^2}{\partial a^1} \right)}^{\text{Layer 2}} \cdot \overbrace{\left( \frac{\partial a^1}{\partial z^1} \cdot \frac{\partial z^1}{\partial w^1} \right)}^{\text{Layer 1}} \quad (7.17)$$

We see a clear **pattern** here! In fact, this is the procedure we'll use for a neural network with **any** number of layers.

**Concept 9**

We can get all of our **weight gradients** by repeatedly appending to the **chain rule**.

If we want to get the **weight gradient** of layer  $\ell$ , we **terminate** with

$$\underbrace{\frac{\partial a^\ell}{\partial z^\ell}}_{\text{Within layer}} \cdot \underbrace{\frac{\partial z^\ell}{\partial w^\ell}}_{\text{Get weight grad}}$$

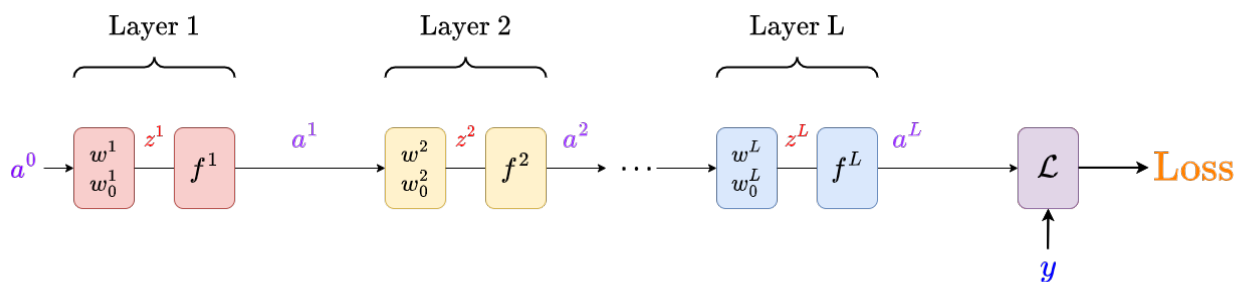
If we want to **extend** to the previous layer, we **instead** multiply by

$$\underbrace{\frac{\partial a^\ell}{\partial z^\ell}}_{\text{Within layer}} \cdot \underbrace{\frac{\partial z^\ell}{\partial a^{\ell-1}}}_{\text{Link layers}}$$

### 7.5.10 Many layers: Doing back-propagation

Now, we'll consider the case of many possible layers.

To make it more readable, we'll use boxes instead of circles for units.



This may look intimidating, but we already have all the tools we need to handle this problem.

Our goal is to get a **gradient** for each of our **weight** vectors  $w^\ell$ , so we can do gradient descent and **improve** our model.

According to our above analysis in Concept 9, we need only a few steps to get all of our gradients.

**Concept 10**

In order to do **back-propagation**, we have to build up our **chain rule** for each weight gradient.

- We start our chain rule with one term shared by every gradient:

$$\overbrace{\frac{\partial \mathcal{L}}{\partial \mathbf{a}^L}}^{\text{Loss unit}}$$

Then, we follow these two steps until we run out of layers:

- We're at layer  $\ell$ . We want to get the **weight gradient** for this layer. We get this by **multiplying** our chain rule by

$$\overbrace{\frac{\partial \mathbf{a}^\ell}{\partial \mathbf{z}^\ell}}^{\text{Within layer}} \cdot \overbrace{\frac{\partial \mathbf{z}^\ell}{\partial \mathbf{w}^\ell}}^{\text{Get weight grad}}$$

We **exclude** this term for any other gradients we want.

- If we aren't at layer 1, there's a previous layer we want to get the weight for. We reach layer  $\ell - 1$  by multiplying our chain rule by

$$\overbrace{\frac{\partial \mathbf{a}^\ell}{\partial \mathbf{z}^\ell}}^{\text{Within layer}} \cdot \overbrace{\frac{\partial \mathbf{z}^\ell}{\partial \mathbf{a}^{\ell-1}}}^{\text{Link layers}}$$

Once we reach layer 1, we have **every single** weight vector we need! Repeat the process for  $w_0$  gradients and then do **gradient descent**.

Let's get an idea of what this looks like in general:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}^\ell} = \overbrace{\left( \frac{\partial \mathcal{L}}{\partial \mathbf{a}^L} \right)}^{\text{Loss unit}} \cdot \overbrace{\left( \frac{\partial \mathbf{a}^L}{\partial \mathbf{z}^L} \cdot \frac{\partial \mathbf{z}^L}{\partial \mathbf{a}^{L-1}} \right)}^{\text{Layer L}} \cdot \overbrace{\left( \frac{\partial \mathbf{a}^{L-1}}{\partial \mathbf{z}^{L-1}} \cdot \frac{\partial \mathbf{z}^{L-1}}{\partial \mathbf{a}^{L-2}} \right)}^{\text{Layer L-1}} \cdot \left( \cdots \right) \cdot \overbrace{\left( \frac{\partial \mathbf{a}^\ell}{\partial \mathbf{z}^\ell} \cdot \frac{\partial \mathbf{z}^\ell}{\partial \mathbf{w}^\ell} \right)}^{\text{Layer } \ell} \quad (7.18)$$

That's pretty ugly. If we need to hide the complexity, we can:



**Notation 11**

If you need to do so for **ease**, you can **compress** your derivatives. For example, if we want to only have the last weight term **separate**, we can do:

$$\frac{\partial \mathcal{L}}{\partial w^\ell} = \overbrace{\frac{\partial \mathcal{L}}{\partial z^\ell}}^{\text{Other}} \cdot \overbrace{\frac{\partial z^\ell}{\partial w^\ell}}^{\text{Weight term}}$$

But we should also explore what each of these terms *are*.

### 7.5.11 What do these derivatives equal?

Let's look at each of these derivatives and see if we can't simplify them a bit.

First, every gradient needs

- The **loss derivative**:

$$\frac{\partial \mathcal{L}}{\partial a^\ell} \quad (7.19)$$

This **depends** on our loss function, so we're **stuck** with that one.

Next, within each layer, we have

- The **activation function** - between our activation  $a$  and preactivation  $z$ :

$$\frac{\partial a^\ell}{\partial z^\ell} \quad (7.20)$$

What does the function between these **look** like?

$$a = f(z) \quad (7.21)$$

Well, that's not super interesting: we **don't know** our function. But, at least we can **write** it using  $f$ : that way, we know that this term only depends on our **activation function**.

$$\frac{\partial a^\ell}{\partial z^\ell} = \overbrace{\left(f^\ell\right)'}^{\text{deriv of func for layer } \ell} \overbrace{\left(z^\ell\right)}^{\text{Deriv input}} \quad (7.22)$$

This expression is a bit visually clunky, but it works. Without the annotation:

$$\frac{\partial a^\ell}{\partial z^\ell} = \left(f^\ell\right)'(z^\ell) \quad (7.23)$$

$z^\ell$  is not being multiplied by  $(f^\ell)'$ , it's the input to that derivative.

Between layers, we have

- We can also think about the derivative of the **linear function** that **connects two layers**:

$$\frac{\partial z^\ell}{\partial a^{\ell-1}} \quad (7.24)$$

So, we want the function of these two:

$$z^\ell = w^\ell a^{\ell-1} + w_0^\ell \quad (7.25)$$

This one is pretty simple! We just take the derivative manually:

$$\frac{\partial z^\ell}{\partial a^{\ell-1}} = w^\ell \quad (7.26)$$

Finally, every gradient will end with...

- The derivative that directly connects to a **weight**, again using the **linear function**:

$$\frac{\partial z^\ell}{\partial w^\ell} \quad (7.27)$$

The linear function is the same:

$$z^\ell = w^\ell a^{\ell-1} + w_0^\ell \quad (7.28)$$

But with a different **variable**, the **derivative** comes out different:

$$\frac{\partial z^\ell}{\partial w^\ell} = a^{\ell-1} \quad (7.29)$$

Be careful not to get this mixed up with the last one! They look similar, but one is within the layer, and the other is between layers.

**Notation 12**

Our **derivatives** for the **chain rule** in a **1-D neural network** take the form:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{a}^L}$$

$$\frac{\partial \mathbf{a}^\ell}{\partial \mathbf{z}^\ell} = (f^\ell)'(\mathbf{z}^\ell)$$


$$\frac{\partial \mathbf{z}^\ell}{\partial \mathbf{a}^{\ell-1}} = \mathbf{w}^\ell$$

$$\frac{\partial \mathbf{z}^\ell}{\partial \mathbf{w}^\ell} = \mathbf{a}^{\ell-1}$$

Now, we can rewrite our generalized expression for gradient:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}^\ell} = \overbrace{\left( \frac{\partial \mathcal{L}}{\partial \mathbf{a}^L} \right)}^{\text{Loss unit}} \cdot \overbrace{\left( (f^L)'(\mathbf{z}^L) \cdot \mathbf{w}^L \right)}^{\text{Layer L}} \cdot \overbrace{\left( (f^{L-1})'(\mathbf{z}^{L-1}) \cdot \mathbf{w}^{L-1} \right)}^{\text{Layer L-1}} \cdot \left( \dots \right) \cdot \overbrace{\left( (f^\ell)'(\mathbf{z}^\ell) \cdot \mathbf{a}^{\ell-1} \right)}^{\text{Layer } \ell} \quad (7.30)$$

Our expressions are more concrete now. It's still pretty visually messy, though.



## 7.5.12 Activation Derivatives

We weren't able to **simplify** our expressions above, partly because we didn't know which **loss** or **activation** function we were going to use.

So, here, we will look at the **common** choices for these functions, and **catalog** what their derivatives look like.

- **Step function**  $\text{step}(z)$ :

$$\frac{d}{dz} \text{step}(z) = 0 \quad (7.31)$$

This is part of why we don't use this function: it has no gradient. We can show this by looking piecewise:

$$\text{step}(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases} \quad (7.32)$$

And take the derivative of each piece:

$$\frac{d}{dz} \text{ReLU}(z) = 0 = \begin{cases} 0 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases} \quad (7.33)$$

- **Rectified Linear Unit**  $\text{ReLU}(z)$ :

$$\frac{d}{dz} \text{ReLU}(z) = \text{step}(z) \quad (7.34)$$

This one might be a bit surprising at first, but it makes sense if you **also** break it up into cases:

$$\text{ReLU}(z) = \max(0, z) = \begin{cases} z & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases} \quad (7.35)$$

And take the derivative of each piece:

$$\frac{d}{dz} \text{ReLU}(z) = \text{step}(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases} \quad (7.36)$$

- **Sigmoid** function  $\sigma(z)$ :

$$\frac{d}{dz} \sigma(z) = \sigma(z)(1 - \sigma(z)) = \frac{e^{-z}}{(1 + e^{-z})^2} \quad (7.37)$$

This derivative is useful for simplifying NLL, and has a nice form.

As a reminder, the function looks like:

We can just compute the derivative with the single-variable chain rule.

$$\sigma(z) = \frac{1}{1 + e^{-z}} \quad (7.38)$$

- **Identity** ("linear") function  $f(z) = z$ :

$$\frac{d}{dz} z = 1 \quad (7.39)$$

This one follows from the definition of the derivative.

We cannot rely on a linear activation function for our **hidden** layers, because a linear neural network is no more **expressive** than one layer.

But, we use it as the output activation for **regression**.

- **Softmax** function  $\text{softmax}(z)$ :

This function has a difficult derivative we won't go over here.

If you're curious, here's a [link](#).

- **Hyperbolic tangent** function  $\tanh(z)$ :

$$\frac{d}{dz} \tanh(z) = 1 - \tanh(z)^2 \quad (7.40)$$

This strange little expression is 1 minus the "hyperbolic secant" squared. We won't bother further with it.

### Notation 13

For our various **activation** functions, we have the **derivatives**:

Step:

$$\frac{d}{dz} \text{step}(z) = 0$$

ReLU:

$$\frac{d}{dz} \text{ReLU}(z) = \text{step}(z)$$

Sigmoid:

$$\frac{d}{dz} \sigma(z) = \sigma(z)(1 - \sigma(z))$$

Identity/Linear:

$$\frac{d}{dz} z = 1$$

## 7.5.13 Loss derivatives

Now, we look at the loss derivatives.

- **Square loss** function  $\mathcal{L}_{sq} = (a - y)^2$ :

$$\frac{d}{da} \mathcal{L}_{sq} = 2(a - y) \quad (7.41)$$

Follows from chain rule+power rule, used for regression.

- **Linear loss** function  $\mathcal{L}_{sq} = |a - y|$ :

$$\frac{d}{da} \mathcal{L}_{lin} = \text{sign}(a - y) \quad (7.42)$$

This one can also be handled piecewise, like  $\text{step}(z)$  and  $\text{ReLU}(z)$ :

$$|u| = \begin{cases} u & \text{if } z \geq 0 \\ -u & \text{if } z < 0 \end{cases} \quad (7.43)$$

We take the piecewise derivative:

$$\frac{d}{du}|u| = \text{sign}(u) = \begin{cases} 1 & \text{if } z \geq 0 \\ -1 & \text{if } z < 0 \end{cases} \quad (7.44)$$

- **NLL** (Negative-Log Likelihood) function  $\mathcal{L}_{\text{NLL}} = -(y \log(a) + (1 - y) \log(1 - a))$

$$\frac{d}{da} \mathcal{L}_{\text{NLL}} = -\left(\frac{y}{a} - \frac{1-y}{1-a}\right) \quad (7.45)$$

- **NLLM** (Negative-Log Likelihood Multiclass) function  $\mathcal{L}_{\text{NLL}} = -\sum_j y_j \log(a_j)$

Similar to softmax, we will omit this derivative.

#### Notation 14

For our various **loss** functions, we have the **derivatives**:

Square:

$$\frac{d}{da} \mathcal{L}_{sq} = 2(a - y)$$

Linear (Absolute):

$$\frac{d}{da} \mathcal{L}_{lin} = \text{sign}(a - y)$$

NLL (Negative-Log Likelihood):

$$\frac{d}{da} \mathcal{L}_{\text{NLL}} = -\left(\frac{y}{a} - \frac{1-y}{1-a}\right)$$

### 7.5.14 Many neurons per layer

Now, we just have left the elephant in the room: what do we do about the case where we have *big* layers? That is, what if we have **multiple** neurons per layer? This makes this more complex.

Well, the solution is the same as earlier in the course: we introduce **matrices**.

But this time, with a twist: we have to do serious **matrix** calculus: a difficult topic indeed.

To handle this, we will go in somewhat **reversed** order, but one that better fits our needs.

- We begin by considering how the chain rule looks when we switch to matrix form.
- We give a general idea of what matrix derivatives look like.
- We list some of the results that matrix calculus gives us, for particular derivatives.
- We actually reason about how matrix calculus *works*.

The last of these is by far the **hardest**, and warrants its own section. Nevertheless, even without it, you can more or less get the idea of what we need - hence why we're going in reversed order.

### 7.5.15 The chain rule: Matrix form

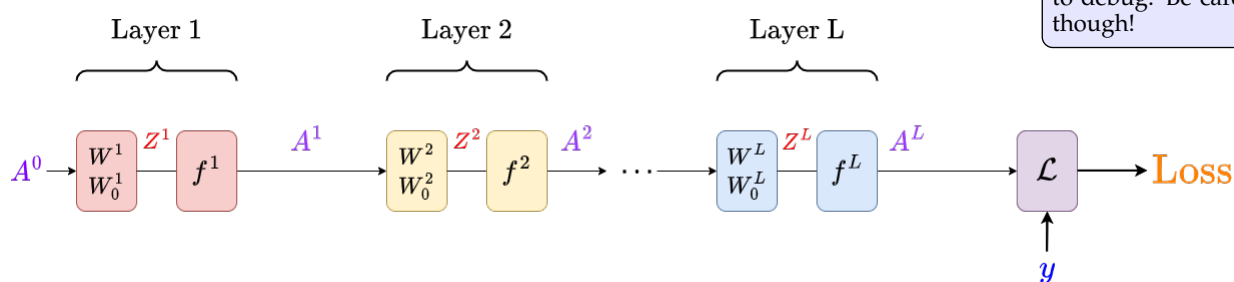
Let's start with the first: the punchline, how does the chain rule and our gradient descent **change** when we add **matrices**?

It turns out, not much: by using **layers** in the last section, we were able to create a pretty powerful and mathematically **tidy** object.

- With layers, each layer feeds into the **next**, with no other interaction. And neurons **within** the same layer do **not** directly **interact** with each other, which simplifies our math greatly.
  - Basically, we have a bunch of functions (neurons) that, within a layer, have **nothing** to do with each other, and only **output** to the **next** layer of similar functions.
- So, we can often **oversimplify** our model by thinking of each **layer** as like a "big" function, taking in a vector of size  $m^\ell$  and outputting a vector of size  $n^\ell$ .

Our main concern is making sure we have agreement of **dimensions**!

So, here's how our model looks now:



In fact, if you just rearranging your matrices and transposing them can be a helpful way to debug. Be careful, though!

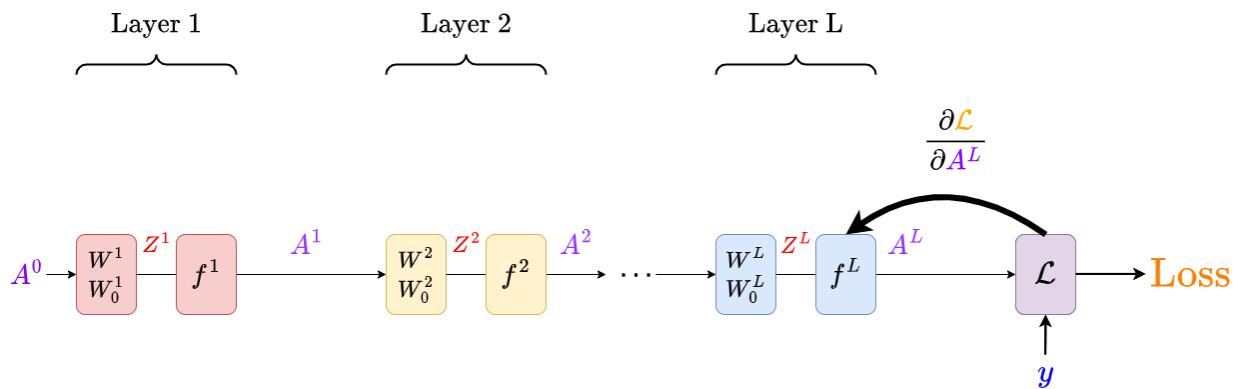
Pretty much the same! Only major difference: swapped scalars for vectors, and vectors for matrices (represented by switching to uppercase)

And, we do backprop the same way, too.

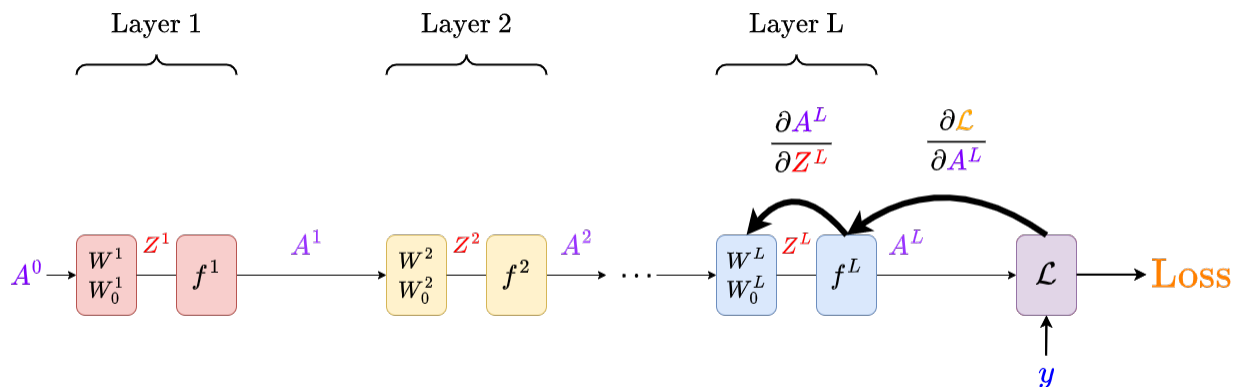
Here, we're not going to explain much as we go: all we're doing is getting the **derivatives** we need for our **chain rule**!

As we go **backwards**, we can build the gradient for each **weight** we come across, in the way we described above.

As always, we start from the loss function:



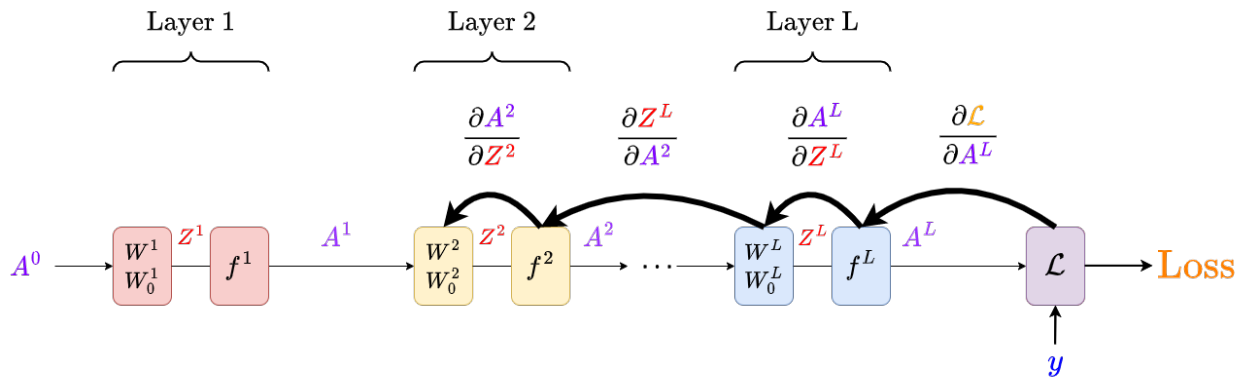
Take another step:



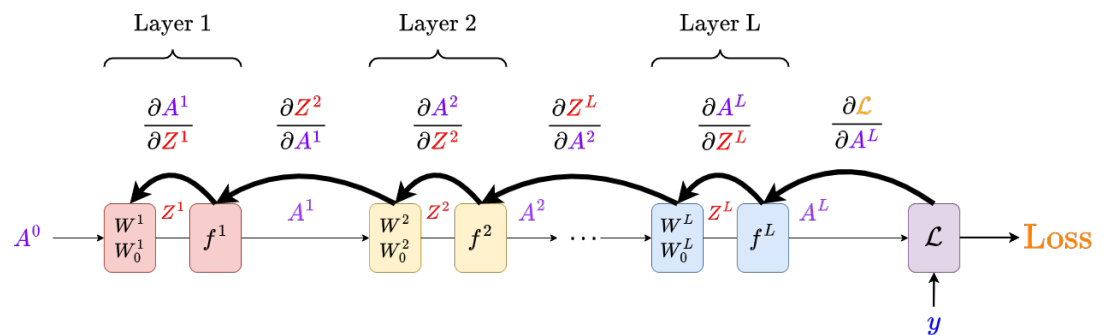
We'll pick up the pace: we'll jump to layer 2 and get its gradient.

The term  $\partial Z^L / \partial A^2$  contains lots of derivatives from every layer between L and 2. But, all we're omitting is the same kinds of steps we're doing in layers 1, 2, and L.





Now, we finally get to layer 1!



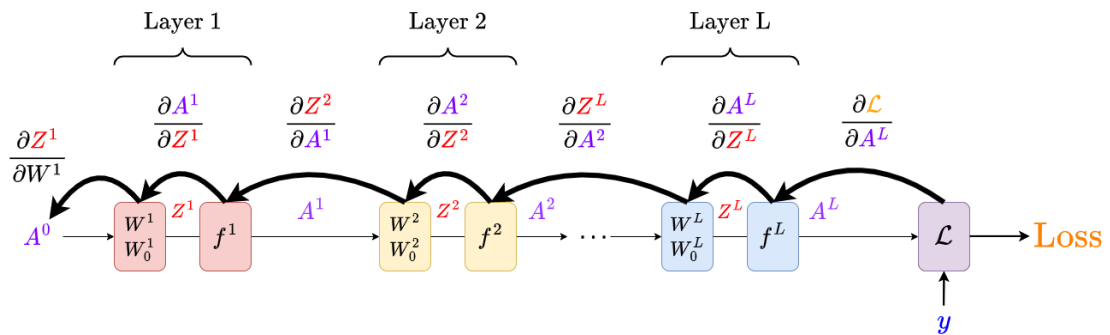
We finish off by getting what we're after: the gradient for  $W^1$ .

#### Notation 15

We depict neural network gradient descent using the below diagram (outside the box):

The **right-facing straight** arrows come **first**: they're part of the **forward pass**, where we get all of our values.

The **left-facing curved** arrows come **after**: they represent the **back-propagation** of the gradient.



And, with this, we can rewrite our general equation for neural network gradients.

### 7.5.16 How the Chain Rule changes in Matrix form

As we discussed before, we can't just add onto our weight gradient to reach another layer: the final term

$$\frac{\partial Z^\ell}{\partial W^\ell} \quad (7.46)$$

Ends our chain rule when we add it:  $W^\ell$  isn't part of the input or output, so it doesn't connect to the previous layer.

So, for this section, we'll add it **separately** at the end of our chain rule:

$$\frac{\partial \mathcal{L}}{\partial W^\ell} = \overbrace{\frac{\partial Z^\ell}{\partial W^\ell}}^{\text{Weight link}} \cdot \overbrace{\left( \frac{\partial \mathcal{L}}{\partial Z^\ell} \right)^T}^{\text{Other layers}}$$

That way, we can add onto  $\partial \mathcal{L} / \partial Z^\ell$  without worrying about the weight derivative.

Notice two minor changes caused by the switch to matrices:

- The order has to be **reversed**.
- We also have to do some weird **transposing**.

Both of these mostly boil down to trying to be careful about **shape**/dimension agreement.

There are also deeper interpretations, but they aren't worth digging into for now.

**Notation 16**


The **gradient**  $\nabla_{W^\ell} \mathcal{L}$  for a neural network is given as:

$$\frac{\partial \mathcal{L}}{\partial W^\ell} = \overbrace{\frac{\partial \mathcal{L}}{\partial Z^\ell}}^{\text{Weight link}} \cdot \overbrace{\left( \frac{\partial \mathcal{L}}{\partial Z^\ell} \right)^T}^{\text{Other layers}}$$

We get our remaining terms  $\partial \mathcal{L} / \partial Z^\ell$  by our usual chain rule:

$$\frac{\partial \mathcal{L}}{\partial Z^\ell} = \overbrace{\left( \frac{\partial \mathcal{L}}{\partial Z^\ell} \right)}^{\text{Layer } \ell} \cdot \left( \cdots \right) \cdot \overbrace{\left( \frac{\partial \mathcal{L}}{\partial Z^{L-1}} \cdot \frac{\partial \mathcal{L}}{\partial A^{L-1}} \right)}^{\text{Layer } L-1} \cdot \overbrace{\left( \frac{\partial \mathcal{L}}{\partial Z^L} \cdot \frac{\partial \mathcal{L}}{\partial A^L} \right)}^{\text{Layer } L} \cdot \overbrace{\left( \frac{\partial \mathcal{L}}{\partial A^L} \right)}^{\text{Loss unit}}$$

This is likely our most important equation in this chapter!

**7.5.17 Relevant Derivatives**

If you aren't interested in understanding matrix derivatives, here we provide the general format of each of the derivatives we care about.

**Notation 17**

Here, we give useful **derivatives** for **neural network gradient descent**.

Loss is not given, so we can't compute it, as before:

$$\overbrace{\frac{\partial \mathcal{L}}{\partial \mathbf{A}^L}}^{(n^L \times 1)}$$

We get the same result for each of these terms as we did before, except in matrix form.

$$\overbrace{\frac{\partial \mathbf{Z}^\ell}{\partial \mathbf{W}^\ell}}^{(m^\ell \times 1)} = \mathbf{A}^{\ell-1}$$

$$\overbrace{\frac{\partial \mathbf{Z}^\ell}{\partial \mathbf{A}^{\ell-1}}}^{(m^\ell \times n^\ell)} = \mathbf{W}^\ell$$

The last one is actually pretty different from before:

$$\overbrace{\frac{\partial \mathbf{A}^\ell}{\partial \mathbf{Z}^\ell}}^{(n^\ell \times n^\ell)} = \begin{bmatrix} f'(z_1^\ell) & 0 & 0 & \dots & 0 \\ 0 & f'(z_2^\ell) & 0 & \dots & 0 \\ 0 & 0 & f'(z_3^\ell) & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & 0 & f'(z_r^\ell) \end{bmatrix}$$

Where  $r$  is the length of  $\mathbf{Z}^\ell$ .

- In short, we only have the  $z_i$  derivative on the  $i^{\text{th}}$  diagonal
- Why? Check the **matrix derivative explanatory notes**.

**Example:** Suppose you have the activation  $f(z) = z^2$ .

Your pre-activation might be

$$\mathbf{z}^\ell = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} \quad (7.47)$$

The output would be

For  $\partial \mathbf{Z}^\ell / \partial \mathbf{W}^\ell$ , check section A.9.2.

For  $\partial \mathbf{Z}^\ell / \partial \mathbf{A}^{\ell-1}$ , check section A.9.3.

For  $\partial \mathbf{a}^\ell / \partial \mathbf{z}^\ell$ , check section A.9.4.

$$\mathbf{a}^\ell = f(\mathbf{z}^\ell) = \begin{bmatrix} 1 \\ 2^2 \\ 3^2 \end{bmatrix} \quad (7.48)$$

But the derivative would be:

$$f(\mathbf{z}) = 2z \quad (7.49)$$

Which, gives our matrix derivative as:

$$\frac{\partial \mathbf{a}^\ell}{\partial \mathbf{z}^\ell} = \begin{bmatrix} f'(1) & 0 & 0 \\ 0 & f'(2) & 0 \\ 0 & 0 & f'(3) \end{bmatrix} = \begin{bmatrix} 2 \cdot 1 & 0 & 0 \\ 0 & 2 \cdot 2 & 0 \\ 0 & 0 & 2 \cdot 3 \end{bmatrix} = \begin{bmatrix} 2 & 0 & 0 \\ 0 & 4 & 0 \\ 0 & 0 & 6 \end{bmatrix}$$

If you want to be able to **derive** some of the derivatives, without reading the matrix derivative section, just use this formula for vector derivatives:

If you have time, do read – you won't understand what you're doing otherwise!

$$\frac{\partial \mathbf{w}}{\partial \mathbf{v}} = \begin{bmatrix} \frac{\partial w_1}{\partial v_1} & \frac{\partial w_2}{\partial v_1} & \dots & \frac{\partial w_n}{\partial v_1} \\ \frac{\partial w_1}{\partial v_2} & \frac{\partial w_2}{\partial v_2} & \dots & \frac{\partial w_n}{\partial v_2} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial w_1}{\partial v_m} & \frac{\partial w_2}{\partial v_m} & \dots & \frac{\partial w_n}{\partial v_m} \end{bmatrix} \quad \left. \begin{array}{l} \text{Column } j \text{ matches } w_j \\ \text{Row } i \text{ matches } v_i \end{array} \right\} \quad (7.50)$$

We can use this for scalars as well: we just treat them as a vector of length 1.

With some cleverness, you can derive the Scalar/Matrix and Matrix/Scalar derivatives as well.

This is contained in the matrix derivatives chapter.

#### Clarification 18

Note that we have chosen a **convention** for how our matrices work: plenty of other resources use a transposed version of matrix derivatives.

This alternate version means the exact **same** thing as our version. Our choice is called the **denominator layout notation** for matrix derivatives.

## 7.6 Training

### 7.6.1 Comments

A few important side notes on training. First, on derivatives:

#### Concept 19

Sometimes, depending on your **loss** and **activation** function, it may be easier to directly compute

$$\frac{\partial \mathcal{L}}{\partial \mathbf{Z}^L}$$

Than it is to find

$$\partial \mathcal{L} / \partial \mathbf{A}^L \text{ and } \partial \mathbf{A}^L / \partial \mathbf{Z}^L$$

So, our algorithm may change slightly.

Another thought: initialization.

#### Concept 20

We typically try to pick a **random initialization**. This does two things:

- Allows us to avoid weird **numerical** and **symmetry** issues that happen when we start with  $\mathbf{W}_{ij} = 0$ .
- We can hopefully find different **local minima** if we run our algorithm multiple times.
  - This is also helped by picking **random data points** in **SGD** (our typical algorithm).

Here, we choose our **initialization** from a **Gaussian** distribution, if you know what that is.

If you do not know a gaussian distribution, that shouldn't be a problem. It is also known as a "normal" distribution.

### 7.6.2 Pseudocode

Our training algorithm for backprop can follow smoothly from what we've laid out.

Here, we'll use the @ symbol to indicate matrix multiplication, following numpy conventions.

SGD-NEURAL-NET( $\mathcal{D}_n, T, L, (m^1, \dots, m^L), (f^1, \dots, f^L), \text{Loss}$ )

```

1  for every layer:
2      Randomly initialize
3          the weights in every layer
4          the biases in every layer
5
6  While termination condition not met:
7      Get random data point i
8      Keep track of time t
9
10     Do forward pass
11         for every layer:
12             Use previous layer's output: get pre-activation
13             Use pre-activation: get new output, activation
14
15         Get loss: forward pass complete
16
17     Do back-propagation
18         for every layer in reversed order:
19             If final layer: #Loss function
20                 Get  $\partial \mathcal{L} / \partial A^L$ 
21
22             Else:
23                 Get  $\partial \mathcal{L} / \partial A^\ell$ : #Link two layers
24                      $(\partial Z^{\ell+1} / \partial A^\ell) @ (\partial \mathcal{L} / \partial Z^{\ell+1})$ 
25
26                 Get  $\partial \mathcal{L} / \partial Z^\ell$ : #Within layer
27                      $(\partial A^\ell / \partial Z^\ell) @ (\partial \mathcal{L} / \partial A^\ell)$ 
28
29             Compute weight gradients:
30                 Get  $\partial \mathcal{L} / \partial W^\ell$ : #Weights
31                      $\partial Z^\ell / \partial W^\ell = A^{\ell-1}$ 
32                      $(\partial Z^\ell / \partial W^\ell) @ (\partial \mathcal{L} / \partial Z^\ell)$ 
33
34                 Get  $\partial \mathcal{L} / \partial W_0^\ell$ : #Biases
35                      $\partial \mathcal{L} / \partial W_0^\ell = (\partial \mathcal{L} / \partial Z^\ell)$ 
36
37             Follow Stochastic Gradient Descent (SGD): #Take step
38                 Update weights:
39                      $W^\ell = W^\ell - (\eta(t) * (\partial \mathcal{L} / \partial W^\ell))$ 
40
41                 Update biases:
42                      $W_0^\ell = W_0^\ell - (\eta(t) * (\partial \mathcal{L} / \partial W_0^\ell))$ 
43
44  Return final neural network with weights and biases
```

*Last Updated: 12/19/23 15:38:18*

## Terms

- Forward pass
- Back-Propagation
- Weight gradient
- Matrix Derivative
- Partial Derivative
- Multivariable Chain Rule
- Total Derivative
- Size of a matrix
- Planar Approximation
- Scalar/scalar derivative
- Vector/scalar derivative
- Scalar/vector derivative
- Vector/vector derivative