

Explanatory Notes for 6.390

Shauntclair Ruiz (Current TA)

Fall 2022


What's still missing?

Last chapter, we used our linear regression model to do classification: we created a "hyper-plane" to **separate** the the data that we placed in each class.

We also mentioned that regularization can increase **structural error**, by limiting what possible θ models we're allowed to use.

But, what if our linear model is already **too limited**? What if we need a more complicated model? This is true in a lot of real-world problems, like vibration:

Our goal was to decrease estimation error, but that's beside the point right now.



This wave doesn't seem particular friendly to a planar approximation.

These kinds of situations are called, appropriately, **non-linear**.

Concept 1

Non-linear behavior cannot be accurately represented by any **linear** model.

In order to create an accurate model, we have to use some **nonlinear** operation.


If we could create effective, non-linear models, we might even be able to deal with data that was previously "**linearly inseparable**".



Possible Solutions: Polynomials

Let's try to think of ways to approach this problem. We'll start with a 1-D input, for simplicity.

Upon hearing "non-linear", we might remember the function we introduced last chapter: the **sigmoid**.



Your friendly neighborhood sigmoid.

Can we use this to create a new model class? For now, unfortunately not: remember that we used this in the last chapter, and we still got a **linear** separator. The reasons were discussed there.

Instead, we can get inspiration from our example of "structural error". For now, let's focus on **regression** (though classification isn't too different):

We'll show ways we can use this kind of approach, when we discuss Neural Networks.




Figure 1: A linear function can't represent this dataset. However, a parabola can!

We're still using our input variable x , but this time, we've "**transformed**" it: we have squared x , giving us a model of the form

Remember that x is 1-D right now!

$$h(x) = Ax^2 + Bx + C \quad (1)$$

It should be clear that this model is more **expressive** than the one before: it can create every model that our linear approach could (just by setting $A = 0$), and it can create new models in a parabola shape.

Reminder: "expressiveness" or "richness" of a hypothesis class is how many models it can represent: a more expressive model can handle more different situations.

Concept 2

We can make our **linear** model more **expressive** by add a squared term, and turning it into a **parabolic** function.

This concept can be extended even further, to any **polynomial**.

Transformation

How do we *generalize* this concept? Well, we have a set of constant parameters A, B, C . These are similar to our constants θ_i . Let's change our notation:

$$h(x) = \theta_2 x^2 + \theta_1 x + \theta_0 \quad (2)$$

Now, we've got something more familiar. We could imagine extending this to any number of terms $\theta_i x^i$: if we needed a cubic function, for example, we could include $\theta_3 x^3$.

This is starting to look pretty similar to our previous model: in fact, we could even separate out θ as a variable:

Notice that θ_0 corresponds to $x^0 = 1$.

$$h(x) = \underbrace{\sum_{i=1}^k \theta_i x^i}_{\text{Polynomial sum}} = \overbrace{\begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_k \end{bmatrix} \cdot \begin{bmatrix} 1 \\ x \\ x^2 \\ \vdots \\ x^k \end{bmatrix}}^{\text{Store as vectors}} = \underbrace{\theta^T}_{\text{Simplify}} \begin{bmatrix} 1 \\ x \\ x^2 \\ \vdots \\ x^k \end{bmatrix} \quad (3)$$

This really *is* starting to look like our linear transformation! That's helpful: we might be able to use the techniques we developed before.

In fact, we can argue that they're **equivalent**: we've just changed what our input vector is. Consider our new input $\phi(x)$:

$$\phi(x) = \begin{bmatrix} 1 \\ x \\ x^2 \\ \vdots \\ x^k \end{bmatrix} \quad h(x) = \theta^T \underbrace{\phi(x)}_{\text{New input}} \quad (4)$$

This is called **transforming** our input.

Definition 3


A **transformation** $\phi(x)$ takes our input vector x and converts it into a **new** vector.

This transformation can be used to:

- Allow our model to handle new, more **complex** situations (more **expressiveness**)
- **Pre-process** our data to make it **easier** for our model to find **patterns**.
- Convert our data into a **usable** format (if, say, the original format doesn't fit into our equations)

Example: Taking our input x and converting it into a polynomial is a **transformation** of our input.

This chapter will focus on these kinds of transformations.



Features

One benefit of only changing out input is that we can continue to use our linear representation: we will be able to optimize a "linear" model θ , over data that has been made **nonlinear**.

These transformations can be complex, especially for multi-dimensional inputs. In this first case, we only combined one input with **itself**. But, often, we can combine multiple together!

Thus, we should be careful to distinguish each input variable from each other. We often call these "**features**". However, we need to be careful:

Clarification 4

We often use the word **feature** in related (but not identical) contexts:

- A **feature** can be one **aspect** of our **original data**: for example, whether or not something is a cat or a dog, or the height of a patient.
- A **feature** can also be one mathematical **variable** in our **transformed input**. x_i is a feature of the data, while each variable of $\phi(x)$ is a feature of the transformed data.

Just like how we have an input space, we call the collection of possible values for our features the **feature space**.

Example: x in our previous example was a feature of the data, while x^i is a feature of our transformed vector.

Combined, this is why we called this technique the **feature transformation**: we apply some *transform* to the *features* of a data, to create a new set of *features*.

Since these transforms only apply to our features, we can keep the structure of a linear function:

Definition 5

Feature transformation allows us to do **linear** regression or classification on a set of **features** we have **non-linearly transformed**:

$$h(x) = \theta^T \phi(x)$$

$\phi(x)$ is our (often nonlinear) transformation of our features x .

