

Explanatory Notes for 6.390

Shauntclair Ruiz (Current TA)


Spring 2023

Example of Activation Functions

So, let's look at some possible **activation** functions:


- **Identity** function z :

$$f(z) = z \quad (1)$$



- This function is called an **identity** function because it "preserves the identity" of the input: the output is the same.
- This is an example of a **linear** function.
 - * As we described in the last section*, linear activation can't make our model more **expressive**.
 - * So, we **almost never** use it (or any other **linear** function) as an activation for a **hidden** layer.
- We mainly use this as an **output** activation function: it allows our final output to be any real number.
 - * This is a good activation function for a **regression** model, which returns a **real** number.
 - * It's a simple function, that can return **any** real number. By contrast, sigmoid and ReLU both have **limited** output ranges.
- **Step** function $\text{step}(z)$:

$$\text{step}(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases} \quad (2)$$




- This function is basically a **sign** function, but uses $\{0, 1\}$ instead of $\{-1, +1\}$.
- Step functions were a common early choice, but because they have a **zero** gradient, we can't use **gradient descent**, and so we basically **never** use them.

Same reason we replaced the sign function with sigmoid.

- **Rectified Linear Unit** $\text{ReLU}(z)$:

$$\text{ReLU}(z) = \max(0, z) = \begin{cases} z & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases} \quad (3)$$




- This is a very **common** choice for activation function, even though the derivative is undefined at 0.
- We specifically use it for internal ("**hidden**") layers: layers that are neither the **first** nor **last** layer.

They're "hidden" because they aren't visible to the input or output.

- **Sigmoid** function $\sigma(z)$:

$$\sigma(z) = \frac{1}{1 + e^{-z}} \quad (4)$$




- This is the **activation** function for our LLC neuron from before.
- Just like LLC, it's useful for the **output neuron** in **binary classification**.
- Can be interpreted as the **probability** of a positive (+1) binary classification.
- We can also use this for multiclass when classes are **NOT** disjoint: we use one sigmoid per class.

* Each sigmoid tells us how likely the data point is to be in that class.

- **Hyperbolic Tangent** $\tanh(z)$:

$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}} \quad (5)$$



- This is function looks similar to sigmoid over a different **range**.
- Unfortunately, it will not get much use in this class.

- **Softmax** function $\text{softmax}(z)$:

$$\text{softmax}(z) = \begin{bmatrix} \exp(z_1) / \sum_i \exp(z_i) \\ \vdots \\ \exp(z_n) / \sum_i \exp(z_i) \end{bmatrix} \quad (6)$$

- Behaves a like a **multi-class** version of **sigmoid**.
- Appropriately, we use it as the **output neuron** for **multi-class** classification.

- Can be interpreted as the **probability** of our k possible classifications.
- * "Disjoint" probability: each option is separate. Sum of the rows adds up to 1.

Concept 1

For the different **activation functions**:

- $f(z) = z$ isn't used for **hidden** layers, but we can use it for regression **output**.
- $\text{sign}(z)$ is **rarely** used.
- $\text{ReLU}(z)$ is often used for "**hidden**" layers.
- $\sigma(z)$ is often used as the **output** for **binary classification**.
- $\text{softmax}(z)$ is often used as the **output** for **multi-class classification**

$\tanh(z)$ is useful, but not a focus of this class.

Remember this caveat, though:

Clarification 2

Multi-class depends on whether a **data point** can be in **multiple classes at the same time**.

- $\text{softmax}(z)$ assumes our classes are **disjoint**: you can only be in **one** class.
 - This is usually what people mean by **multi-class**.
- $\sigma(z)$ can be used when classes are **not disjoint**: you can be in **multiple** classes.
 - You can think of this as **binary classification** for each class.

When using sigmoids, we need **one** sigmoid for each **class**.

Example: We can compare use cases for each of these:

- Softmax could be used to answer, "which word is the next one in the sentence?"
 - Every word in a sentence is only followed by one word: they're mutually exclusive.
- Sigmoids could be used to answer, "what genre of book is that?"
 - A book is often in more than one genre.

Loss functions and activation functions

As we can see above, your **activation** function depends on what kind of **problem** you're dealing with.

The same is true for our **loss** function: we used **different** loss functions for classification and regression.

Classification can be further broken up into **binary** versus **multiclass** classification.

To summarize our findings, we'll **sort** this information:

Concept 3

Each of our **tasks** requires a different **loss** and output **activation** function.


We emphasize that we specifically mean the **output** activation function: the activation function used in **hidden layers** doesn't have to match the loss function.

task	f^L	Loss
Regression	Linear z	Squared $(g - y)^2$
Binary Class	Sigmoid $\sigma(z)$	NLL $y \log g + (1 - y) \log(1 - g)$
Multi-Class	Softmax $\text{softmax}(z)$	NLLM $\sum_j y_j \log(g_j)$

Special Case: If we allow **multiple** classes at the **same** time (non-disjoint), we use **binary** classification for each of them, rather than multi-class.

Example: An example for each type:

- **Regression:** Predicting the amount of rainfall in centimeters tomorrow.
- **Binary Classification:** Will the stock market go up or down tomorrow?
- **Multi-Class:** What species of tree is this?
- **Multiple Binary:** What are the themes in this movie?



Other Considerations

You might consider using other functions, based on the needs of a more specialized task. We'll ignore those cases, for the most part.

But, if you want to try a new function, the **data type** is the most important for whether we can use it.

Concept 4

If you want to use a new **activation** or **loss** function, you have to pay attention to the **input/output** type.

Example: $\tanh(z)$ outputs over the range $(-1, 1)$. We could use it, if that was the range we wanted.

Be careful, though:

Clarification 5

It's important to stress that while our **output activation** depends on the task, **hidden layers** don't have to.

Hidden layers can use one of several **different** activation functions, regardless of the **task**.

However, some activation functions tend to be **better** for making a model than others.

Example: Often, we use ReLU for hidden layers, but it's rarely used as an output activation function.

We also might use **sigmoid** as a hidden layer for a regression model, even though regression most commonly uses a **linear** output.